## TP 1 : Reminder on Markov Chains - Stochastic gradient descent

## Exercise 1 : Box-Muller and Marsaglia-Bray algorithm

Let $R$ a random variable with Rayleigh distribution with parameter 1, whose probability density function $f_{R}$ is given below, and $\Theta$ with uniform distribution on $[0,2 \pi]$. We also assume that $R$ and $\Theta$ are independent. We have

$$
\forall r \in \mathbb{R}, \quad f_{R}(r)=r \exp \left(-\frac{r^{2}}{2}\right) \mathbb{1}_{\mathbb{R}^{+}}(r)
$$

1. Let $X$ and $Y$ such that

$$
X=R \cos (\Theta) \quad \text { and } \quad Y=R \sin (\Theta)
$$

Prove that both $X$ and $Y$ have $\mathcal{N}(0,1)$ distribution and are independent.
2. Write an algorithm for sampling 2 independent Gaussian distributions $\mathcal{N}(0,1)$.

```
Algorithm 1: Marsaglia-Bray algorithm
while \(V_{1}^{2}+V_{2}^{2}>1\) do
    Sample \(U_{1}, U_{2}\) independent r.v. with distribution \(\mathcal{U}([0,1])\);
    Set \(V_{1}=2 U_{1}-1\) and \(V_{2}=2 U_{2}-1\).
end
Set \(S=\sqrt{-2 \log \left(V_{1}^{2}+V_{2}^{2}\right)}\);
Set \(X=S \frac{V_{1}}{\sqrt{V_{1}^{2}+V_{2}^{2}}}\) and \(Y=S \frac{V_{2}}{\sqrt{V_{1}^{2}+V_{2}^{2}}}\);
return \((X, Y)\).
```

3. Consider the algorithm given above.
a) What is the distribution of $\left(V_{1}, V_{2}\right)$ at the end of the "while" loop?
b) Set

$$
T_{1}=\frac{V_{1}}{\sqrt{V_{1}^{2}+V_{2}^{2}}}, \quad T_{2}=\frac{V_{2}}{\sqrt{V_{1}^{2}+V_{2}^{2}}} \quad \text { and } \quad V=V_{1}^{2}+V_{2}^{2}
$$

Show that $\left(T_{1}, T_{2}\right)$ and $V$ are independent, $V \sim \mathcal{U}([0,1])$ and $\left(T_{1}, T_{2}\right)$ has the same distribution as $(\cos (\Theta), \sin (\Theta))$ with $\Theta \sim \mathcal{U}([0,2 \pi])$.
c) What is the distribution of the output $(X, Y)$ ?
d) What is the expected number of steps in the "while" loop?

Teaching assistants : P. Clavier (pierre.clavier@polytechnique.edu), M. Noble (maxence.noble-bourillot@polytechnique.edu)

## Exercise 2: Invariant distribution

We define a Markov chain $\left(X_{n}\right)_{n \geq 0}$ with values in $[0,1]$ as follows : given the current value $X_{n}(n \in \mathbb{N})$ of the chain,

- if $X_{n}=\frac{1}{m}$ (for some positive integer $m$ ), we let :

$$
\begin{cases}X_{n+1}=\frac{1}{m+1} & \text { with probability } 1-X_{n}^{2} \\ X_{n+1} \sim \mathcal{U}([0,1]) & \text { with probability } X_{n}^{2}\end{cases}
$$

- if not, $X_{n+1} \sim \mathcal{U}([0,1])$.

1. Prove that the transition kernel of the chain $\left(X_{n}\right)_{n \geq 0}$ is given by :

$$
P(x, A)= \begin{cases}x^{2} \int_{A \cap[0,1]} d t+\left(1-x^{2}\right) \delta_{\frac{1}{m+1}}(A) & \text { if } x=\frac{1}{m} \\ \int_{A \cap[0,1]} d t & \text { otherwise }\end{cases}
$$

where $\delta_{\alpha}$ is the Dirac measure at $\alpha$.
2. Prove that the uniform distribution on $[0,1]$ is invariant for $P$. In the following, this invariant distribution will be denoted by $\pi$.
3. Let $x \notin\left\{\frac{1}{m}, m \in \mathbb{N}^{*}\right\}$. Compute the value of $\operatorname{Pf}(x)=\mathbb{E}\left[f\left(X_{1}\right) \mid X_{0}=x\right]$, for a bounded measurable function $f$. Deduce $P^{n} f(x)$ for all $n \geqslant 1$. Compute $\lim _{n \rightarrow+\infty} P^{n} f(x)$ in terms of $\int f(x) \pi(x) d x$.
4. Let $x=\frac{1}{m}$ with $m \geqslant 2$.
a) Let $n \in \mathbb{N}^{*}$. Compute $P^{n}\left(x, \frac{1}{n+m}\right)$ in terms of $m$ and $n$.
b) Do we have $\lim _{n \rightarrow+\infty} P^{n}(x, A)=\pi(A)$ when $A=\bigcup_{q \in \mathbb{N}}\left\{\frac{1}{m+1+q}\right\}$ ?

[^0]
## Exercise 3 : Stochastic Gradient Learning in Neural Networks, [Bot91, BCN16]

In the exercise, we consider the problem of classifying patterns $x$ into two classes $y= \pm 1$. We assume that there is a relationship between a pattern and its class, embodied by some probability distribution $\mathbb{P}(x, y)$. If we know this distribution, we know the conditional probabilities $\mathbb{P}(y \mid x)$ as well, and we can solve immediately the problem using the Bayes decision rule. Learning means "Acquiring enough knowledge about $\mathbb{P}(x, y)$ from the examples to solve the classification problem".

The statistical machine learning approach begins with the collection of a sizeable set of examples $\left\{\left(x_{1}, y_{1}\right), \ldots,\left(x_{n}, y_{n}\right)\right\}$, where for each $i \in \llbracket 1, n \rrbracket$ the vector $x_{i}$ represents the features and the scalar $y_{i}$ a label indicating whether $x_{i}$ belongs $\left(y_{i}=1\right)$ or not $\left(y_{i}=-1\right)$ to a particular class. With such a set of examples, one can construct a classification program, defined by a prediction function $h$, and measure its performance by counting how often the program prediction $h\left(x_{i}\right)$ differs from the correct prediction $y_{i}$. To avoid rote memorization, one should aim to find a prediction function that generalizes the concepts that may be learned from the examples. One way to achieve good generalized performance is to choose amongst a carefully selected class of prediction functions.

Thanks to such a high-dimensional sparse representation of documents, it has been deemed empirically sufficient to consider prediction functions of the form $h(x ; w, \tau)=w^{t} x-\tau$. Here, $w^{t} x$ is a linear discriminant parameterized by $w \in \mathbb{R}^{d}$ and $\tau \in \mathbb{R}$ is a bias that provides a way to compromise between precision and recall, $\mathbb{P}[y=1 \mid h(x)=1]$ and $\mathbb{P}[h(x)=1 \mid y=1]$ respectively. The accuracy of the predictions could be determined by counting the number of times that $\operatorname{sign}(h(x ; w, \tau))$ matches the correct label, i.e., 1 or -1 . However, while such a prediction function may be appropriate for classifying new features, formulating an optimization problem around it to choose the parameters $(w ; \tau)$ is impractical in large-scale settings due to the combinatorial structure introduced by the sign function, which is discontinuous. Instead, one typically employs a continuous approximation through a loss function that measures a cost for predicting $h$ when the true label is $y$.

An Adaline (Widrow and Hoff, 1960) actually learns by (i) considering linear prediction function, $h(x, w)=w^{t} x$, and (ii) measuring the quality of the system through the mean squared error :

$$
C_{\text {Adaline }}(w)=\int(y-h(x, w))^{2} \mathrm{~d} \mathbb{P}(x, y)=\int\left(y-w^{t} x\right)^{2} \mathrm{~d} \mathbb{P}(x, y)
$$

Learning consists of finding the parameter $w^{\star}$ that minimizes the above, or a more general, cost. This framework is the basis of classical statistical inference theory. Hundreds of practical algorithms have been derived.

[^1]In the following, we will denote by $z=(x, y)$ the observation and consider the cost or expected risk given a parameter vector $w$ with respect to the probability $\mathbb{P}$

$$
R(w)=\mathbb{E}[J(w, z)]=\int\left(y-w^{t} x\right)^{2} \mathrm{~d} \mathbb{P}(z)
$$

While it may be desirable to minimize the expected loss that would be incurred from any inputoutput pair, such a goal is untenable when one does not have complete information about $\mathbb{P}$. Thus, in practice, one seeks the solution of a problem that involves an estimate of the expected risk $R$. In supervised learning, one has access (either all-at-once or incrementally) to a set of $n \in N$ independently drawn input-output samples $\left\{z_{i}=\left(x_{i}, y_{i}\right)\right\}_{i=1}^{n}$ and one may define the empirical risk function $R_{n}: \mathbb{R}^{d} \rightarrow \mathbb{R}_{+}$by

$$
R_{n}(w)=\frac{1}{n} \sum_{i=1}^{n}\left(y_{i}-w^{t} x_{i}\right)^{2}
$$

1. Describe the stochastic gradient descent algorithm for minimizing the empirical risk and implement it.
2. Sample a set of observations $\left\{z_{i}\right\}_{i=1}^{n}$ by generating a collection of random points $x_{i}$ of $\mathbb{R}^{2}, \bar{w} \in \mathbb{R}^{2}$ seen as the normal vector of an hyperplane, a straight line here, and assigning the label $y_{i}$ according to the side of the hyperplane where the point $x_{i}$ is.
3. Test the algorithm you wrote at the first question over these observations. What is the vector $w^{\star}$ estimated? Is it far from $\bar{w}$ ?

4. Noise your observations $\left\{z_{i}\right\}_{i=1}^{n}$ with an additive Gaussian noise and perform the optimisation again. Compare with the result of question 3.
5. Test the algorithm on the Breast Cancer Wisconsin (Diagnostic) Data Set [WSM95] : http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Diagnostic\).

## Références

[BCN16] Léon Bottou, Frank E. Curtis, and Jorge Nocedal. Optimization methods for largescale machine learning. eprint arXiv :1606.04838, 2016.
[Bot91] Léon Bottou. Stochastic gradient learning in neural networks. In Neuro-N̂̂mes 91, 1991.
[WSM95] William H. Wolberg, W. Nick Street, and Olvi L. Mangasarian. UCI machine learning repository, 1995.

[^2]
[^0]:    Teaching assistants : P. Clavier (pierre.clavier@polytechnique.edu), M. Noble (maxence.noble-bourillot@polytechnique.edu)

    2/4
    Send your work at compstatsmva@gmail.com.

[^1]:    Teaching assistants : P. Clavier (pierre.clavier@polytechnique.edu), M. Noble (maxence.noble-bourillot@polytechnique.edu)

[^2]:    Teaching assistants : P. Clavier (pierre.clavier@polytechnique.edu), M. Noble (maxence.noble-bourillot@polytechnique.edu)

